

Starbucks_Capstone_Project

December 22, 2019

1 Starbucks Capstone Challenge

1.1 1. Definition

1.1.1 Introduction

This data set contains simulated data that mimics customer behavior on the Starbucks rewards mobile app. Once every few days, Starbucks sends out an offer to users of the mobile app. An offer can be merely an advertisement for a drink or an actual offer such as a discount or BOGO (buy one get one free). Some users might not receive any offer during certain weeks.

Not all users receive the same offer, and that is the challenge to solve with this data set.

Your task is to combine transaction, demographic and offer data to determine which demographic groups respond best to which offer type. This data set is a simplified version of the real Starbucks app because the underlying simulator only has one product whereas Starbucks actually sells dozens of products.

Every offer has a validity period before the offer expires. As an example, a BOGO offer might be valid for only 5 days. You'll see in the data set that informational offers have a validity period even though these ads are merely providing information about a product; for example, if an informational offer has 7 days of validity, you can assume the customer is feeling the influence of the offer for 7 days after receiving the advertisement.

You'll be given transactional data showing user purchases made on the app including the timestamp of purchase and the amount of money spent on a purchase. This transactional data also has a record for each offer that a user receives as well as a record for when a user actually views the offer. There are also records for when a user completes an offer.

Keep in mind as well that someone using the app might make a purchase through the app without having received an offer or seen an offer.

1.1.2 Example

To give an example, a user could receive a discount offer buy 10 dollars get 2 off on Monday. The offer is valid for 10 days from receipt. If the customer accumulates at least 10 dollars in purchases during the validity period, the customer completes the offer.

However, there are a few things to watch out for in this data set. Customers do not opt into the offers that they receive; in other words, a user can receive an offer, never actually view the offer, and still complete the offer. For example, a user might receive the "buy 10 dollars get 2 dollars off offer", but the user never opens the offer during the 10 day validity period. The customer spends

15 dollars during those ten days. There will be an offer completion record in the data set; however, the customer was not influenced by the offer because the customer never viewed the offer.

1.1.3 Cleaning

This makes data cleaning especially important and tricky.

You'll also want to take into account that some demographic groups will make purchases even if they don't receive an offer. From a business perspective, if a customer is going to make a 10 dollar purchase without an offer anyway, you wouldn't want to send a buy 10 dollars get 2 dollars off offer. You'll want to try to assess what a certain demographic group will buy when not receiving any offers.

1.1.4 Final Advice

Because this is a capstone project, you are free to analyze the data any way you see fit. For example, you could build a machine learning model that predicts how much someone will spend based on demographics and offer type. Or you could build a model that predicts whether or not someone will respond to an offer. Or, you don't need to build a machine learning model at all. You could develop a set of heuristics that determine what offer you should send to each customer (i.e., 75 percent of women customers who were 35 years old responded to offer A vs 40 percent from the same demographic to offer B, so send offer A).

1.1.5 Data Sets

The data is contained in three files:

- `portfolio.json` - containing offer ids and meta data about each offer (duration, type, etc.)
- `profile.json` - demographic data for each customer
- `transcript.json` - records for transactions, offers received, offers viewed, and offers completed

Here is the schema and explanation of each variable in the files:

portfolio.json * `id` (string) - offer id * `offer_type` (string) - type of offer ie BOGO, discount, informational * `difficulty` (int) - minimum required spend to complete an offer * `reward` (int) - reward given for completing an offer * `duration` (int) - time for offer to be open, in days * `channels` (list of strings)

profile.json * `age` (int) - age of the customer * `became_member_on` (int) - date when customer created an app account * `gender` (str) - gender of the customer (note some entries contain 'O' for other rather than M or F) * `id` (str) - customer id * `income` (float) - customer's income

transcript.json * `event` (str) - record description (ie transaction, offer received, offer viewed, etc.) * `person` (str) - customer id * `time` (int) - time in hours since start of test. The data begins at time `t=0` * `value` - (dict of strings) - either an offer id or transaction amount depending on the record

Note: If you are using the workspace, you will need to go to the terminal and run the command `conda update pandas` before reading in the files. This is because the version of pandas in the workspace cannot read in the `transcript.json` file correctly, but the newest version of pandas can. You can access the terminal from the orange icon in the top left of this notebook.

1.1.6 Problem Statement

The problem is simple: we want to make better purchasing offers to Starbucks' customers. For this, we can use customer's past behaviour to find patterns and try to be more assertive. As given by the Udacity's Starbucks Project Overview, the basic task is to use the data to identify which groups of people are most responsive to each type of offer, and how best to present each type of offer. In other words, this is a classification problem where the model takes user behaviour data as input and produces a group as output (either previously defined or not).

This has been one of the [most used](#) applications of machine learning in the industry, since it provides you with means to save money spent on marketing campaigns by directing content to users who are more likely to convert based on a multitude of characteristics.

1.1.7 Metrics

To evaluate the trained models, we'll compare the models based on it's F1-Score. This is a widely used metric to evaluate classification problems. [Aditya Mishra defines](#) it as follows: > F1 Score is the Harmonic Mean between precision and recall. The range for F1 Score is [0, 1]. It tells you how precise your classifier is (how many instances it classifies correctly), as well as how robust it is (it does not miss a significant number of instances).

Since the F1-Score is [a weighted average of the precision and recall metrics](#), it works well to evaluate datasets that aren't very well balanced.

```
[1]: import pandas as pd
import numpy as np
import math
import json
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn import metrics
from sklearn.preprocessing import MultiLabelBinarizer, LabelEncoder
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.dummy import DummyClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn import svm
from sklearn.tree import DecisionTreeClassifier
from datetime import datetime

%matplotlib inline

[2]: # read in the json files
portfolio = pd.read_json('data/portfolio.json', orient = 'records', lines = _
    ↪ True)
profile = pd.read_json('data/profile.json', orient = 'records', lines = True)
transcript = pd.read_json('data/transcript.json', orient = 'records', lines = _
    ↪ True)
```

1.1.8 portfolio dataset preparation

```
[3]: portfolio.shape
```

```
[3]: (10, 6)
```

```
[4]: portfolio.dtypes
```

```
[4]: channels      object
difficulty      int64
duration        int64
id              object
offer_type      object
reward          int64
dtype: object
```

```
[5]: portfolio.head(10)
```

```
[5]:
```

	channels	difficulty	duration	\
0	[email, mobile, social]	10	7	
1	[web, email, mobile, social]	10	5	
2	[web, email, mobile]	0	4	
3	[web, email, mobile]	5	7	
4	[web, email]	20	10	
5	[web, email, mobile, social]	7	7	
6	[web, email, mobile, social]	10	10	
7	[email, mobile, social]	0	3	
8	[web, email, mobile, social]	5	5	
9	[web, email, mobile]	10	7	

	id	offer_type	reward
0	ae264e3637204a6fb9bb56bc8210ddfd	bogo	10
1	4d5c57ea9a6940dd891ad53e9dbe8da0	bogo	10
2	3f207df678b143eea3cee63160fa8bed	informational	0
3	9b98b8c7a33c4b65b9aebfe6a799e6d9	bogo	5
4	0b1e1539f2cc45b7b9fa7c272da2e1d7	discount	5
5	2298d6c36e964ae4a3e7e9706d1fb8c2	discount	3
6	fafdc668e3743c1bb46111dcafc2a4	discount	2
7	5a8bc65990b245e5a138643cd4eb9837	informational	0
8	f19421c1d4aa40978ebb69ca19b0e20d	bogo	5
9	2906b810c7d4411798c6938adc9daaa5	discount	2

We can see some variables that could use some preprocessing, such as `channels` which is a list of values. We'll use [Scikit-learn's MultiLabelBinarizer](#) to transform it.

Then we'll create a copy of the original `portfolio` dataset with the new encoded columns.

```
[6]: mlb = MultiLabelBinarizer()
      encoded_channels = mlb.fit_transform(portfolio['channels'])
```

```
[7]: encoded_channels_df = pd.DataFrame(encoded_channels,
                                         columns = mlb.classes_,
                                         index = portfolio['channels'].index)
```

```
[8]: processed_portfolio = portfolio.copy()

      processed_portfolio = processed_portfolio.join(encoded_channels_df)
```

```
[9]: processed_portfolio = processed_portfolio.drop('channels', axis = 1)
      print(processed_portfolio)
```

	difficulty	duration	id	offer_type \
0	10	7	ae264e3637204a6fb9bb56bc8210ddfd	bogo
1	10	5	4d5c57ea9a6940dd891ad53e9dbe8da0	bogo
2	0	4	3f207df678b143eea3cee63160fa8bed	informational
3	5	7	9b98b8c7a33c4b65b9aebfe6a799e6d9	bogo
4	20	10	0b1e1539f2cc45b7b9fa7c272da2e1d7	discount
5	7	7	2298d6c36e964ae4a3e7e9706d1fb8c2	discount
6	10	10	fafdc668e3743c1bb461111dcafc2a4	discount
7	0	3	5a8bc65990b245e5a138643cd4eb9837	informational
8	5	5	f19421c1d4aa40978ebb69ca19b0e20d	bogo
9	10	7	2906b810c7d4411798c6938adc9daaa5	discount

	reward	email	mobile	social	web
0	10	1	1	1	0
1	10	1	1	1	1
2	0	1	1	0	1
3	5	1	1	0	1
4	5	1	0	0	1
5	3	1	1	1	1
6	2	1	1	1	1
7	0	1	1	1	0
8	5	1	1	1	1
9	2	1	1	0	1

```
[10]: processed_portfolio.isna().sum()
```

```
[10]: difficulty    0
      duration      0
      id            0
      offer_type    0
      reward        0
      email         0
      mobile        0
```

```
social      0
web         0
dtype: int64
```

No empty cells, looks like this dataset is good enough for now.

1.1.9 profile dataset preparation

```
[11]: profile.shape
```

```
[11]: (17000, 5)
```

```
[12]: profile.dtypes
```

```
[12]: age                int64
became_member_on      int64
gender                object
id                    object
income                float64
dtype: object
```

```
[13]: profile.head(10)
```

```
[13]:
```

	age	became_member_on	gender	id	income
0	118	20170212	None	68be06ca386d4c31939f3a4f0e3dd783	NaN
1	55	20170715	F	0610b486422d4921ae7d2bf64640c50b	112000.0
2	118	20180712	None	38fe809add3b4fcf9315a9694bb96ff5	NaN
3	75	20170509	F	78afa995795e4d85b5d9ceeca43f5fef	100000.0
4	118	20170804	None	a03223e636434f42ac4c3df47e8bac43	NaN
5	68	20180426	M	e2127556f4f64592b11af22de27a7932	70000.0
6	118	20170925	None	8ec6ce2a7e7949b1bf142def7d0e0586	NaN
7	118	20171002	None	68617ca6246f4fbc85e91a2a49552598	NaN
8	65	20180209	M	389bc3fa690240e798340f5a15918d5c	53000.0
9	118	20161122	None	8974fc5686fe429db53ddde067b88302	NaN

The `profile` dataset has some interesting characteristics, such as: * `NaN` values for the `income` column. * `None` value for the `gender` column (which is ok, since there are more genders than M, F, but this might need to be treated since we have `0` value - that probably means 'others'). * `became_member_on` has dates that are treated as numbers. * `age` has the value `118` repeated several times. Since it's fairly unlikely for the customers to be 118 years-old, I'm assuming this is also some kind of previous data preprocessing and won't change it.

```
[14]: profile['gender'].unique()
```

```
[14]: array([None, 'F', 'M', '0'], dtype=object)
```

```
[15]: profile.isna().sum()
```

```
[15]: age                0
      became_member_on    0
      gender             2175
      id                 0
      income             2175
      dtype: int64
```

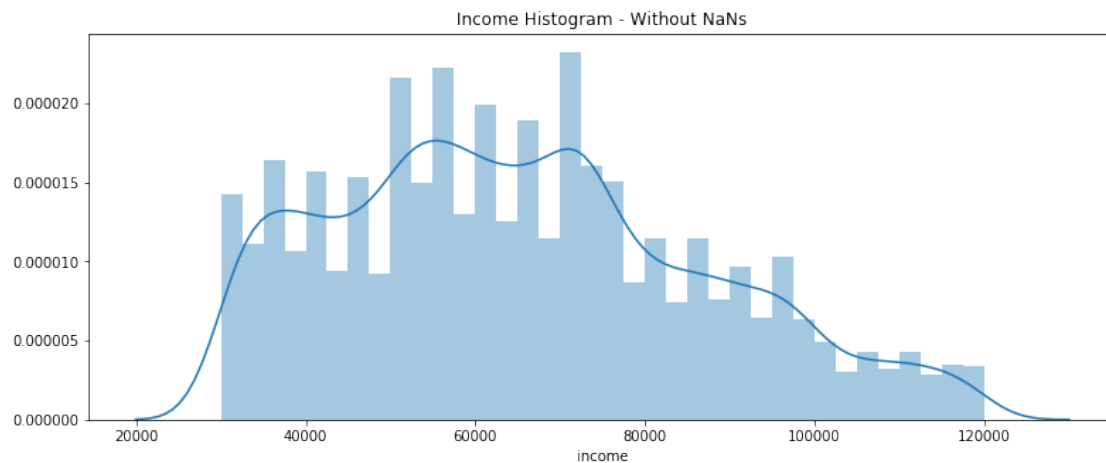
The `None` value for `gender` is not understood as a value, so it needs treatment. We'll fill the `None` values with `Unknown`.

```
[16]: processed_profile = profile.copy()

      processed_profile['gender'].fillna('Unknown', inplace = True)
```

For the `income` column, let's see what's the data distribution without the `NaNs`:

```
[17]: plt.figure(figsize = (13, 5))
      sns.distplot(profile['income'].dropna());
      plt.title('Income Histogram - Without NaNs');
```



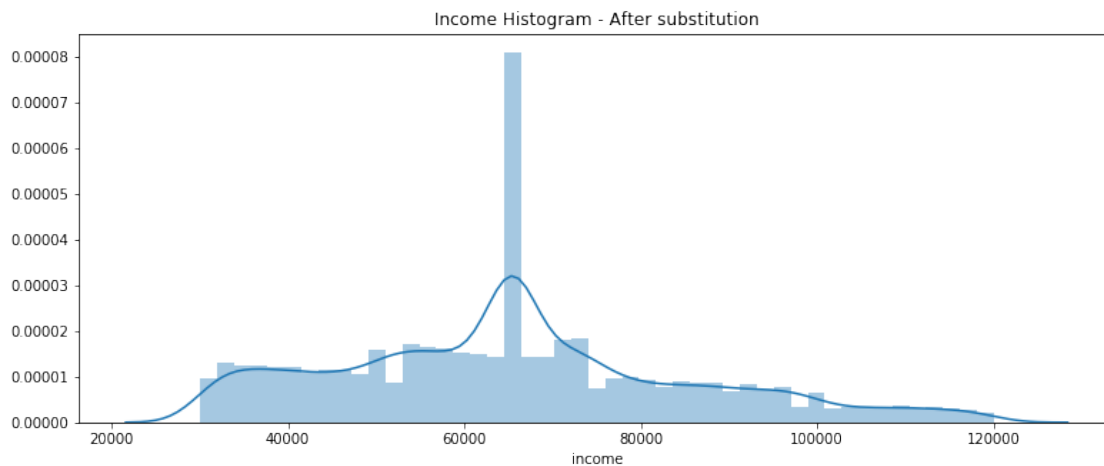
```
[18]: profile['income'].describe()
```

```
[18]: count      14825.000000
      mean       65404.991568
      std       21598.299410
      min       30000.000000
      25%      49000.000000
      50%      64000.000000
      75%      80000.000000
      max      120000.000000
      Name: income, dtype: float64
```

By filling the NaNs with some value, we'll change the distribution's shape. Since both mean and median are fairly close to each other, it doesn't matter much which one we'll use. I chose the mean.

```
[19]: processed_profile['income'].fillna(profile['income'].mean(), inplace = True)
```

```
[20]: plt.figure(figsize = (13, 5))
sns.distplot(processed_profile['income']);
plt.title('Income Histogram - After substitution');
```



We can see it changed a bit the distribution of the data, but we were already expecting that.

```
[21]: processed_profile.isna().sum()
```

```
[21]: age                0
became_member_on       0
gender                 0
id                     0
income                 0
dtype: int64
```

No empty cells, looks like this dataset is good enough for now.

Another alteration we can work on is to transform the `became_member_on` column into a `datetime` object column, which allows us to work with time windows. This means that we can get the information for how long a customer has been a member of the program. Since we don't really know when those campaigns were build, we don't have a specific point in time to measure a customer's "age" as a customer. On the other hand, we can simply use the year the user signed up for the membership.

```
[22]: def convert_to_date(date_element):
        return datetime.strptime(str(date_element), '%Y%m%d')
```



```
[23]: became_member_on_date = processed_profile['became_member_on'].  
      ↪ apply(convert_to_date)
```

```
[24]: became_member_on_date.head()
```

```
[24]: 0    2017-02-12  
      1    2017-07-15  
      2    2018-07-12  
      3    2017-05-09  
      4    2017-08-04  
      Name: became_member_on, dtype: datetime64[ns]
```

```
[25]: became_member_on_year = became_member_on_date.apply(lambda date_point: ↪  
      ↪ date_point.year)
```

```
[26]: processed_profile['became_member_on_year'] = became_member_on_year
```

```
[27]: processed_profile = processed_profile.drop(['became_member_on'], axis = 1)
```

```
[28]: processed_profile.head()
```

```
[28]:   age  gender          id      income \  
0  118  Unknown  68be06ca386d4c31939f3a4f0e3dd783  65404.991568  
1   55         F  0610b486422d4921ae7d2bf64640c50b  112000.000000  
2  118  Unknown  38fe809add3b4fcf9315a9694bb96ff5  65404.991568  
3   75         F  78afa995795e4d85b5d9ceeca43f5fef  100000.000000  
4  118  Unknown  a03223e636434f42ac4c3df47e8bac43  65404.991568  
  
      became_member_on_year  
0                2017  
1                2017  
2                2018  
3                2017  
4                2017
```

1.1.10 transcript dataset preparation

```
[29]: transcript.shape
```

```
[29]: (306534, 4)
```

```
[30]: transcript.dtypes
```

```
[30]: event      object  
      person    object  
      time      int64  
      value     object
```

dtype: object

```
[31]: transcript.isna().sum()
```

```
[31]: event      0
      person    0
      time      0
      value     0
      dtype: int64
```

```
[32]: transcript.sample(10)
```

```
[32]:
```

	event	person	time \
76581	offer viewed	ffdefcac307f4ca99ac1ebd51470f106	186
30709	transaction	885d8e7de35e4b7abb7c6f494c82c271	48
87734	transaction	25a6a2df56ba4a4a9ad898dbf083c6b5	222
106692	offer viewed	8d20cb72c43c43699ff1074fd1e7d468	306
19710	offer viewed	95dc4202ae22461d865b13da40b3d557	12
103990	transaction	c1dc09fa7f5940a38c2f393f89cb4a50	294
36890	transaction	b3883064435140ce8feac3c1da259948	72
180137	transaction	eae57eab7df940d898b5b32d968ecdc7	438
185351	transaction	3977011340cf47d08a23c4b3d6f6aef6	450
197364	transaction	c0bbc13872474c63a83e8b503bb88f72	486

	value
76581	{'offer id': '4d5c57ea9a6940dd891ad53e9dbe8da0'}
30709	{'amount': 7.03}
87734	{'amount': 13.44}
106692	{'offer id': '9b98b8c7a33c4b65b9aebfe6a799e6d9'}
19710	{'offer id': 'fafdcd668e3743c1bb461111dcafc2a4'}
103990	{'amount': 1.1400000000000001}
36890	{'amount': 4.93}
180137	{'amount': 7.78}
185351	{'amount': 0.66}
197364	{'amount': 35.26}

This dataset has no empty values. The `value` column has a dictionary that has different keys and values. It will be easier to have this column expanded into multiple columns to work with the events.

```
[33]: value_expanded = transcript['value'].apply(pd.Series)
      value_expanded.head()
```

```
[33]:
```

	offer id	amount	offer_id	reward
0	9b98b8c7a33c4b65b9aebfe6a799e6d9	NaN	NaN	NaN
1	0b1e1539f2cc45b7b9fa7c272da2e1d7	NaN	NaN	NaN
2	2906b810c7d4411798c6938adc9daaa5	NaN	NaN	NaN

3	fafdc668e3743c1bb46111dcafc2a4	NaN	NaN	NaN
4	4d5c57ea9a6940dd891ad53e9dbe8da0	NaN	NaN	NaN

Now that we have this matrix, we can include it to the rest of the transcript data.

```
[34]: transcript_expanded = pd.concat([transcript, value_expanded], axis = 1)
```

```
[35]: transcript_expanded.sample(10)
```

```
[35]:
```

	event	person	time \
97082	transaction	98ba686b0d9943279f4ef1ee9480eaff	258
89483	transaction	35faf6c136ac4eff95cb298191d4dcbe	228
263829	offer viewed	eb9e4fd1814b4fba9cdf0eb00f35278	582
58191	offer received	c77f073996314357b57c8765e85d0a01	168
2724	offer received	cd0cb9d1599c47a0bb69c49d19c2117f	0
176516	offer completed	dcadc299e3914526acf3eac4815b7930	426
228763	transaction	777a1d8f93bb4884829eead23fd6fb53	528
201361	transaction	83b26151b9c9487d897a67be32d2eba4	498
169536	transaction	6c0f89d1092545bbb522281d8337dd2b	414
13088	offer completed	1ba58dab3cae4787ad26d8d77ab7380f	0

	value \
97082	{'amount': 6.82}
89483	{'amount': 5.42}
263829	{'offer_id': 'f19421c1d4aa40978ebb69ca19b0e20d'}
58191	{'offer_id': '2298d6c36e964ae4a3e7e9706d1fb8c2'}
2724	{'offer_id': '4d5c57ea9a6940dd891ad53e9dbe8da0'}
176516	{'offer_id': '2298d6c36e964ae4a3e7e9706d1fb8c2...'}
228763	{'amount': 13.97}
201361	{'amount': 11.76}
169536	{'amount': 12.12}
13088	{'offer_id': 'f19421c1d4aa40978ebb69ca19b0e20d...'}

	offer_id	amount \
97082	NaN	6.82
89483	NaN	5.42
263829	f19421c1d4aa40978ebb69ca19b0e20d	NaN
58191	2298d6c36e964ae4a3e7e9706d1fb8c2	NaN
2724	4d5c57ea9a6940dd891ad53e9dbe8da0	NaN
176516	NaN	NaN
228763	NaN	13.97
201361	NaN	11.76
169536	NaN	12.12
13088	NaN	NaN

	offer_id	reward
97082	NaN	NaN

89483		NaN	NaN
263829		NaN	NaN
58191		NaN	NaN
2724		NaN	NaN
176516	2298d6c36e964ae4a3e7e9706d1fb8c2		3.0
228763		NaN	NaN
201361		NaN	NaN
169536		NaN	NaN
13088	f19421c1d4aa40978ebb69ca19b0e20d		5.0

Now we can manipulate this dataset in order to clean it. We will deal with the duplicate `offer id/offer_id` column and remove unnecessary columns.

The `offer id/offer_id` duplicates seem to relate to the same information, but for different events, which means that each line can never have two values for `offer id/offer_id` (only one per line). To treat that, we can create a single column that receives the presented `offer id/offer_id` value when one exists.

```
[36]: transcript_expanded['offer_id_redux'] = np.where(transcript_expanded['offer_id']
↳ isnull() & transcript_expanded['offer_id'].notnull(),
↳ transcript_expanded['offer_id'],
↳ transcript_expanded['offer_id'])
```

Now we can drop unnecessary columns.

```
[37]: transcript_expanded = transcript_expanded.drop(['value',
↳ 'offer id',
↳ 'offer_id'],
↳ axis = 1)
```

```
[38]: transcript_expanded.head()
```

```
[38]:
```

	event	person	time	amount	reward	\
0	offer received	78afa995795e4d85b5d9ceeca43f5fef	0	NaN	NaN	
1	offer received	a03223e636434f42ac4c3df47e8bac43	0	NaN	NaN	
2	offer received	e2127556f4f64592b11af22de27a7932	0	NaN	NaN	
3	offer received	8ec6ce2a7e7949b1bf142def7d0e0586	0	NaN	NaN	
4	offer received	68617ca6246f4fbc85e91a2a49552598	0	NaN	NaN	

	offer_id_redux
0	9b98b8c7a33c4b65b9aebfe6a799e6d9
1	0b1e1539f2cc45b7b9fa7c272da2e1d7
2	2906b810c7d4411798c6938adc9daaa5
3	fafdc668e3743c1bb46111dcafc2a4
4	4d5c57ea9a6940dd891ad53e9dbe8da0

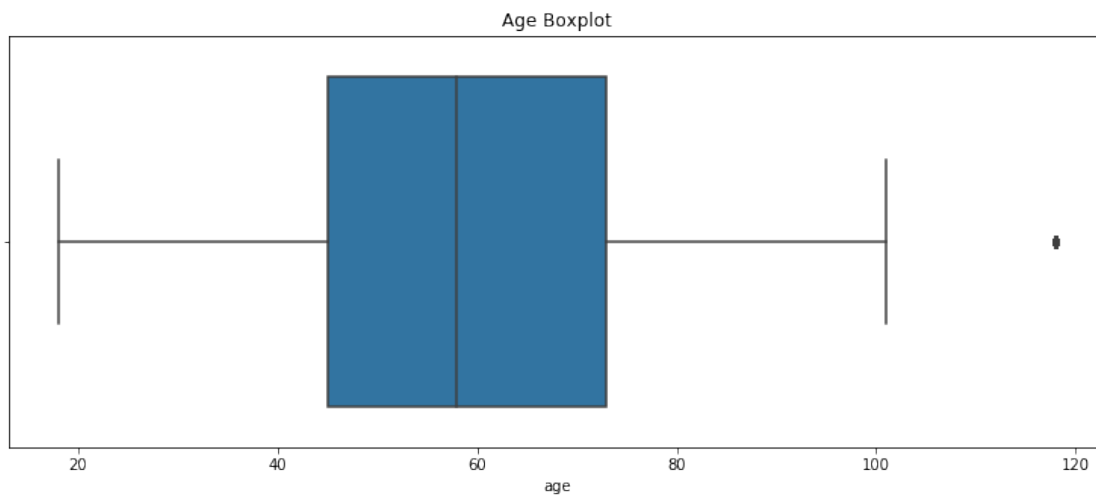
1.2 Exploratory Data Analysis

Starting with the `profile` dataset, we can check a bit of the age and income of the customers.

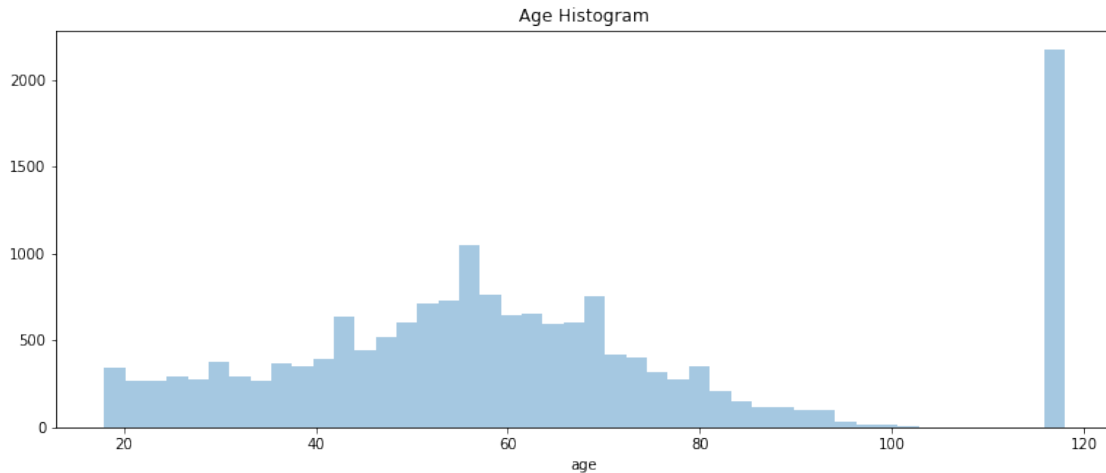
```
[39]: processed_profile['age'].describe()
```

```
[39]: count      17000.000000  
      mean        62.531412  
      std         26.738580  
      min         18.000000  
      25%         45.000000  
      50%         58.000000  
      75%         73.000000  
      max        118.000000  
      Name: age, dtype: float64
```

```
[40]: plt.figure(figsize = (13, 5))  
      sns.boxplot(processed_profile['age']);  
      plt.title('Age Boxplot');
```



```
[41]: plt.figure(figsize = (13, 5))  
      sns.distplot(processed_profile['age'], kde = False);  
      plt.title('Age Histogram');
```

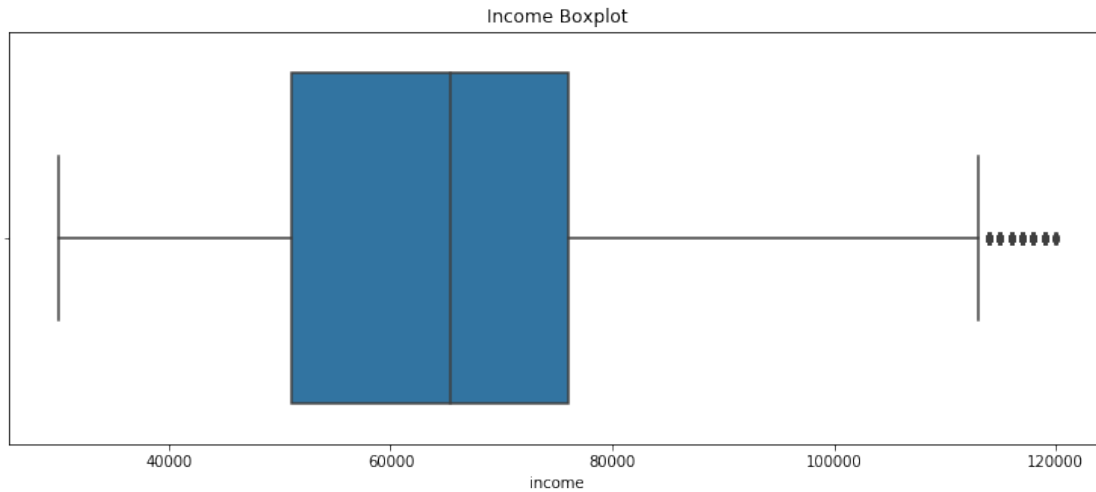


This histogram and the boxplot show us a great amount of outliers near 120. That was expected, since we decided not to treat the 118 values (as discussed earlier). The users are mostly adults, with ages between 40 and 80 and the median in a bit lower than 60 (which is also lower than the mean, that is 62.5).

```
[42]: processed_profile['income'].describe()
```

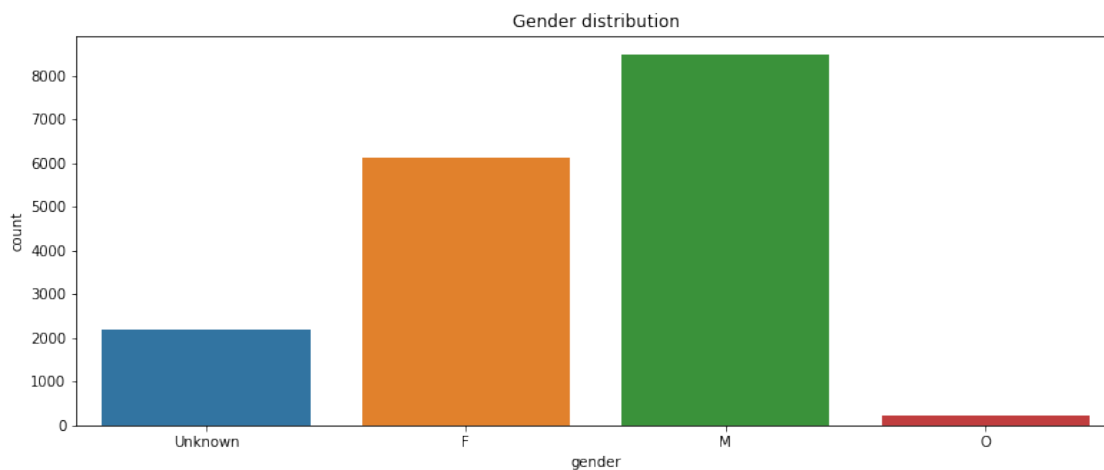
```
[42]: count      17000.000000
      mean       65404.991568
      std        20169.288288
      min        30000.000000
      25%         51000.000000
      50%         65404.991568
      75%         76000.000000
      max        120000.000000
      Name: income, dtype: float64
```

```
[43]: plt.figure(figsize = (13, 5))
      sns.boxplot(processed_profile['income']);
      plt.title('Income Boxplot');
```



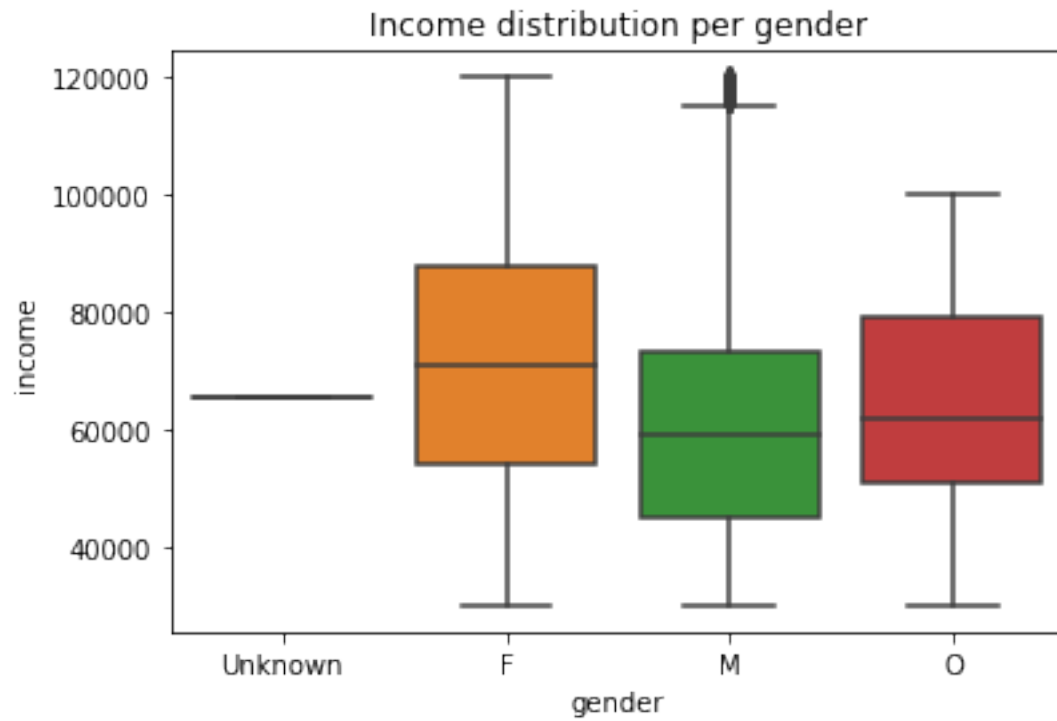
We already saw the histogram for this data, the boxplot shows the distribution is a bit skewed. Mean and median seem close (around \$65,000) and most of the data is between \$60,000 and \$75,000.

```
[44]: plt.figure(figsize = (13, 5))
sns.countplot(processed_profile['gender']);
plt.title('Gender distribution');
```



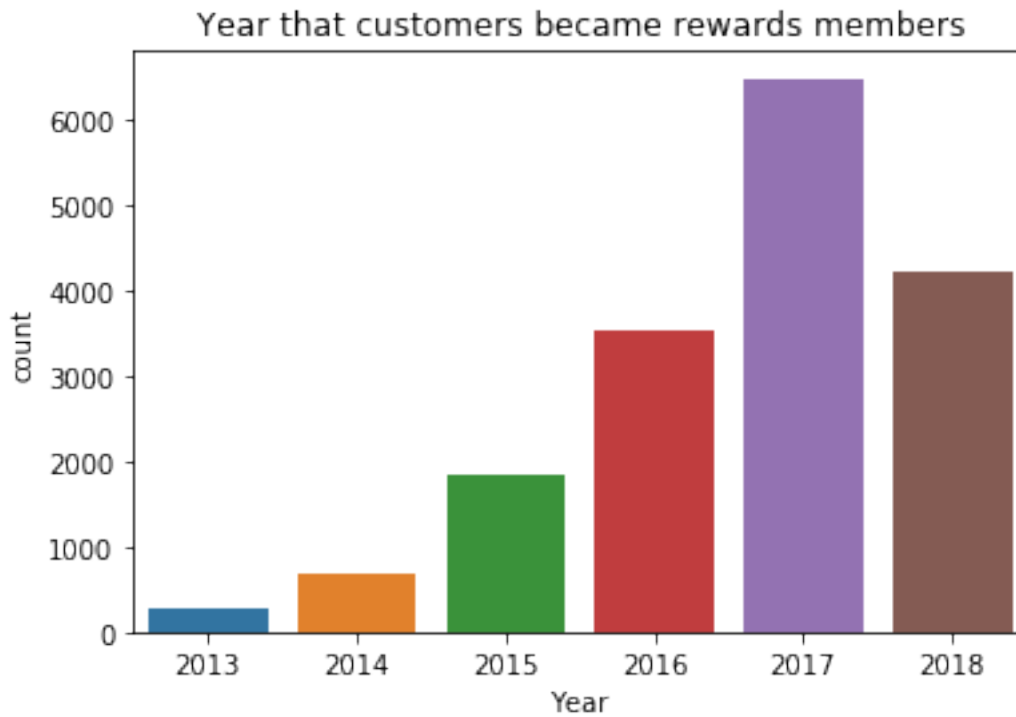
There are more male than females in the dataset. There is a large amount of Unknown and a few 'others'.

```
[45]: sns.boxplot(x = 'gender',
                  y = 'income',
                  data = processed_profile);
plt.title('Income distribution per gender');
```



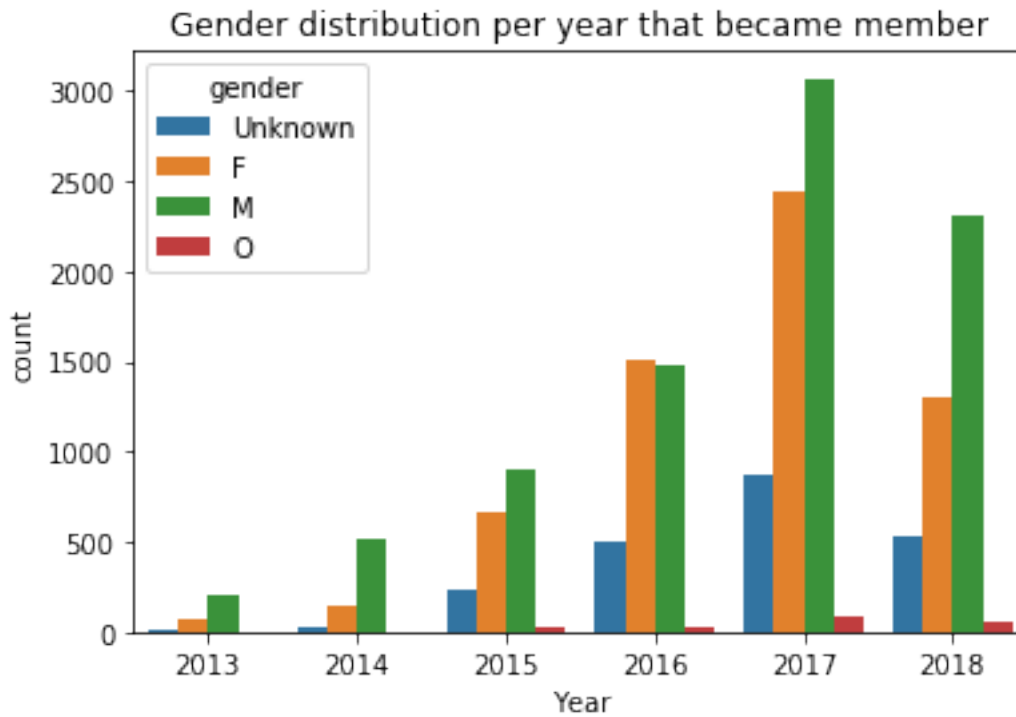
The boxplots show there is no evidence that the income differs accross genders.

```
[46]: sns.countplot('became_member_on_year',  
                  data = processed_profile);  
plt.title('Year that customers became rewards members')  
plt.xlabel('Year');
```

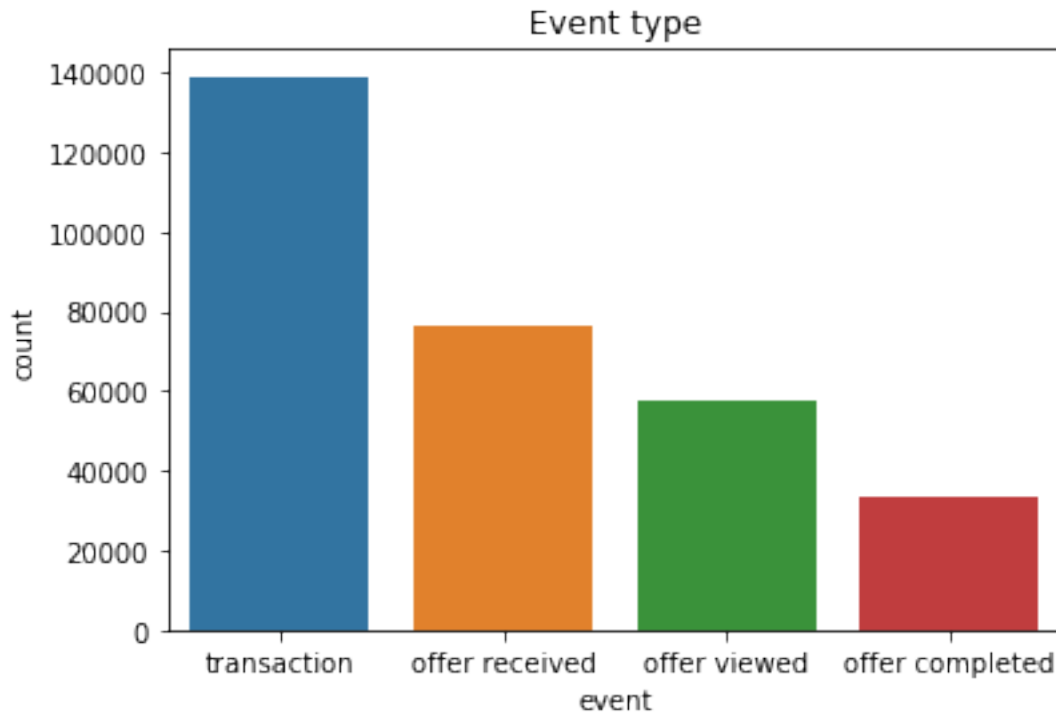
More customers became members in 2017, followed by 2018 then 2016. The first year the with records is 2013 and it's also the year with fewer members signing for the program.

```
[47]: sns.countplot(x = 'became_member_on_year',  
                    hue = 'gender',  
                    data = processed_profile);  
plt.xlabel('Year');  
plt.title('Gender distribution per year that became member');
```



In all years, there are either more male members signing up for the program (with exception of 2016, in which there were slight more female members registrations).

```
[48]: sns.countplot(transcript_expanded['event'],  
                    order = transcript_expanded['event'].value_counts().index;  
plt.title('Event type');
```



The event type that is most frequent is `transaction`, followed by `offer received`, `offer viewed` and `offer completed`. This makes sense, since it looks like a usual customer funnel for marketing.

1.3 Merging datasets and data preprocessing

We can get add more information to the `transcript_expanded` dataset by joining it to the `processed_portfolio` dataset, that contains information on the offers.

```
[49]: transcript_with_portfolio = transcript_expanded.merge(processed_portfolio,
                                                         how = 'left',
                                                         left_on = 'offer_id',
                                                         right_on = 'id')
```

```
[50]: # sort dataset by user and time, to make it easier for humans to see data
transcript_with_portfolio = transcript_with_portfolio.sort_values(['person', 'time'])
```

```
[51]: transcript_with_portfolio.groupby(['event', 'offer_type'])['offer_type'].count()
```

```
[51]: event      offer_type
offer completed  bogo      15669
              discount    17910
offer received  bogo      30499
```

	discount	30543
	informational	15235
offer viewed	bogo	25449
	discount	21445
	informational	10831

Name: offer_type, dtype: int64

BOGO and discount offers are the only ones with an associated `offer completed` event. Also, the `transaction` event doesn't have an associated id, which means that the join won't add data to those.

Basically, we can look into the BOGO and discount offers user funnels in this order: * offer received * offer viewed * transaction * offer completed

For the informational offer, the funnel has this order: * offer received * offer viewed * transaction

This means that we have to find a way to add the transaction data into the funnel. This can be done through the user id and the timestamp. If a user didn't get to the transaction part, it won't have an `offer completed` event as well. For the informational offer, the transaction data will only exist if it came after the `offer viewed` event. The timestamp will help us find the "organic" users: users that became members regardless of the campaigns. All those scenarios can be found by ordering data by time, offer id, person and event.

```
[52]: view_to_complete = transcript_with_portfolio[['time', 'offer_id_redux',
→ 'person', 'event']][(transcript_with_portfolio['event']=='transaction') |
→ (transcript_with_portfolio['event'] == 'offer viewed')].
→ groupby(['person', 'offer_id_redux'])
```

```
[53]: view_to_complete.head()
```

```
[53]:
```

	time	offer_id_redux	\
77705	192	5a8bc65990b245e5a138643cd4eb9837	
89291	228		NaN
139992	372	3f207df678b143eea3cee63160fa8bed	
168412	414		NaN
187554	456	f19421c1d4aa40978ebb69ca19b0e20d	
228422	528		NaN
233413	540	fafdc668e3743c1bb461111dcafc2a4	
237784	552		NaN
258883	576		NaN
85769	216	f19421c1d4aa40978ebb69ca19b0e20d	
284472	630	f19421c1d4aa40978ebb69ca19b0e20d	
16179	6	3f207df678b143eea3cee63160fa8bed	
75427	186	2298d6c36e964ae4a3e7e9706d1fb8c2	
133370	354	5a8bc65990b245e5a138643cd4eb9837	
177937	432	0b1e1539f2cc45b7b9fa7c272da2e1d7	
222679	516	9b98b8c7a33c4b65b9aebfe6a799e6d9	
18431	12	fafdc668e3743c1bb461111dcafc2a4	
174774	426	4d5c57ea9a6940dd891ad53e9dbe8da0	

293372	660	5a8bc65990b245e5a138643cd4eb9837
67584	168	2298d6c36e964ae4a3e7e9706d1fb8c2
131476	348	f19421c1d4aa40978ebb69ca19b0e20d
165442	408	5a8bc65990b245e5a138643cd4eb9837
263808	582	9b98b8c7a33c4b65b9aebfe6a799e6d9
27148	36	5a8bc65990b245e5a138643cd4eb9837
105508	300	fafdc668e3743c1bb461111dcafc2a4
140716	372	3f207df678b143eea3cee63160fa8bed
173080	420	fafdc668e3743c1bb461111dcafc2a4
27263	36	5a8bc65990b245e5a138643cd4eb9837
76441	186	fafdc668e3743c1bb461111dcafc2a4
219332	510	f19421c1d4aa40978ebb69ca19b0e20d
...
26549	36	fafdc668e3743c1bb461111dcafc2a4
72825	180	4d5c57ea9a6940dd891ad53e9dbe8da0
269576	594	f19421c1d4aa40978ebb69ca19b0e20d
103669	288	3f207df678b143eea3cee63160fa8bed
148670	396	4d5c57ea9a6940dd891ad53e9dbe8da0
287565	636	3f207df678b143eea3cee63160fa8bed
68851	168	fafdc668e3743c1bb461111dcafc2a4
132444	348	4d5c57ea9a6940dd891ad53e9dbe8da0
182159	438	f19421c1d4aa40978ebb69ca19b0e20d
221408	510	4d5c57ea9a6940dd891ad53e9dbe8da0
289457	642	ae264e3637204a6fb9bb56bc8210ddfd
16325	6	fafdc668e3743c1bb461111dcafc2a4
83982	210	9b98b8c7a33c4b65b9aebfe6a799e6d9
233587	540	5a8bc65990b245e5a138643cd4eb9837
23028	24	fafdc668e3743c1bb461111dcafc2a4
73167	180	4d5c57ea9a6940dd891ad53e9dbe8da0
168973	414	ae264e3637204a6fb9bb56bc8210ddfd
226096	522	fafdc668e3743c1bb461111dcafc2a4
177513	432	fafdc668e3743c1bb461111dcafc2a4
293268	660	4d5c57ea9a6940dd891ad53e9dbe8da0
15591	6	f19421c1d4aa40978ebb69ca19b0e20d
65899	168	5a8bc65990b245e5a138643cd4eb9837
218451	510	f19421c1d4aa40978ebb69ca19b0e20d
294735	666	9b98b8c7a33c4b65b9aebfe6a799e6d9
15836	6	fafdc668e3743c1bb461111dcafc2a4
69626	174	0b1e1539f2cc45b7b9fa7c272da2e1d7
133074	354	2906b810c7d4411798c6938adc9daaa5
168022	414	2906b810c7d4411798c6938adc9daaa5
230690	534	9b98b8c7a33c4b65b9aebfe6a799e6d9
262475	582	2906b810c7d4411798c6938adc9daaa5

	person	event
77705	0009655768c64bdeb2e877511632db8f	offer viewed
89291	0009655768c64bdeb2e877511632db8f	transaction

139992	0009655768c64bdeb2e877511632db8f	offer viewed
168412	0009655768c64bdeb2e877511632db8f	transaction
187554	0009655768c64bdeb2e877511632db8f	offer viewed
228422	0009655768c64bdeb2e877511632db8f	transaction
233413	0009655768c64bdeb2e877511632db8f	offer viewed
237784	0009655768c64bdeb2e877511632db8f	transaction
258883	0009655768c64bdeb2e877511632db8f	transaction
85769	00116118485d4dfda04fdbaba9a87b5c	offer viewed
284472	00116118485d4dfda04fdbaba9a87b5c	offer viewed
16179	0011e0d4e6b944f998e987f904e8c1e5	offer viewed
75427	0011e0d4e6b944f998e987f904e8c1e5	offer viewed
133370	0011e0d4e6b944f998e987f904e8c1e5	offer viewed
177937	0011e0d4e6b944f998e987f904e8c1e5	offer viewed
222679	0011e0d4e6b944f998e987f904e8c1e5	offer viewed
18431	0020c2b971eb4e9188eac86d93036a77	offer viewed
174774	0020c2b971eb4e9188eac86d93036a77	offer viewed
293372	0020c2b971eb4e9188eac86d93036a77	offer viewed
67584	0020ccbbb6d84e358d3414a3ff76cffd	offer viewed
131476	0020ccbbb6d84e358d3414a3ff76cffd	offer viewed
165442	0020ccbbb6d84e358d3414a3ff76cffd	offer viewed
263808	0020ccbbb6d84e358d3414a3ff76cffd	offer viewed
27148	003d66b6608740288d6cc97a6903f4f0	offer viewed
105508	003d66b6608740288d6cc97a6903f4f0	offer viewed
140716	003d66b6608740288d6cc97a6903f4f0	offer viewed
173080	003d66b6608740288d6cc97a6903f4f0	offer viewed
27263	00426fe3ffde4c6b9cb9ad6d077a13ea	offer viewed
76441	00426fe3ffde4c6b9cb9ad6d077a13ea	offer viewed
219332	004b041fbfe44859945daa2c7f79ee64	offer viewed
...
26549	ffede3b700ac41d6a266fa1ba74b4f16	offer viewed
72825	ffede3b700ac41d6a266fa1ba74b4f16	offer viewed
269576	ffede3b700ac41d6a266fa1ba74b4f16	offer viewed
103669	fff0f0aac6c547b9b263080f09a5586a	offer viewed
148670	fff0f0aac6c547b9b263080f09a5586a	offer viewed
287565	fff0f0aac6c547b9b263080f09a5586a	offer viewed
68851	fff29fb549084123bd046dbc5ceb4faa	offer viewed
132444	fff29fb549084123bd046dbc5ceb4faa	offer viewed
182159	fff29fb549084123bd046dbc5ceb4faa	offer viewed
221408	fff29fb549084123bd046dbc5ceb4faa	offer viewed
289457	fff29fb549084123bd046dbc5ceb4faa	offer viewed
16325	fff3ba4757bd42088c044ca26d73817a	offer viewed
83982	fff3ba4757bd42088c044ca26d73817a	offer viewed
233587	fff3ba4757bd42088c044ca26d73817a	offer viewed
23028	fff7576017104bcc8677a8d63322b5e1	offer viewed
73167	fff7576017104bcc8677a8d63322b5e1	offer viewed
168973	fff7576017104bcc8677a8d63322b5e1	offer viewed
226096	fff7576017104bcc8677a8d63322b5e1	offer viewed

```

177513 fff8957ea8b240a6b5e634b6ee8eafcf offer viewed
293268 fff8957ea8b240a6b5e634b6ee8eafcf offer viewed
15591 fffad4f4828548d1b5583907f2e9906b offer viewed
65899 fffad4f4828548d1b5583907f2e9906b offer viewed
218451 fffad4f4828548d1b5583907f2e9906b offer viewed
294735 fffad4f4828548d1b5583907f2e9906b offer viewed
15836 ffff82501cea40309d5fdd7edcca4a07 offer viewed
69626 ffff82501cea40309d5fdd7edcca4a07 offer viewed
133074 ffff82501cea40309d5fdd7edcca4a07 offer viewed
168022 ffff82501cea40309d5fdd7edcca4a07 offer viewed
230690 ffff82501cea40309d5fdd7edcca4a07 offer viewed
262475 ffff82501cea40309d5fdd7edcca4a07 offer viewed

```

[57730 rows x 4 columns]

Since the dataset is grouped and ordered, we can simply fill the offer id gaps with the previous value.

```
[54]: fill_offer_id = view_to_complete['offer_id_redux'].
      ↪fillna(view_to_complete['offer_id_redux'].ffill()).ffill()
```

Since this column is indexed, we can join it with the original dataset and add the missing values to either one of the redundant columns (we will actually make a new column with the values for sanity check). Then drop the unnecessary columns.

```
[55]: with_full_offer_id = transcript_with_portfolio.join(fill_offer_id,
      ↪rsuffix = '_right')
```

```
[56]: with_full_offer_id['offer_id'] = np.where(with_full_offer_id['offer_id_redux'].
      ↪isnull(),
      ↪
      ↪with_full_offer_id['offer_id_redux_right'],
      ↪with_full_offer_id['offer_id_redux'])
```

```
[57]: with_full_offer_id.head()
```

```
[57]:
```

	event	person	time	amount	\
55972	offer received	0009655768c64bdeb2e877511632db8f	168	NaN	
77705	offer viewed	0009655768c64bdeb2e877511632db8f	192	NaN	
89291	transaction	0009655768c64bdeb2e877511632db8f	228	22.16	
113605	offer received	0009655768c64bdeb2e877511632db8f	336	NaN	
139992	offer viewed	0009655768c64bdeb2e877511632db8f	372	NaN	

	reward_x	offer_id_redux	difficulty	duration	\
55972	NaN	5a8bc65990b245e5a138643cd4eb9837	0.0	3.0	
77705	NaN	5a8bc65990b245e5a138643cd4eb9837	0.0	3.0	
89291	NaN	NaN	NaN	NaN	
113605	NaN	3f207df678b143eea3cee63160fa8bed	0.0	4.0	

139992	NaN	3f207df678b143eea3cee63160fa8bed	0.0	4.0
--------	-----	----------------------------------	-----	-----

	id	offer_type	reward_y	email	\
55972	5a8bc65990b245e5a138643cd4eb9837	informational	0.0	1.0	
77705	5a8bc65990b245e5a138643cd4eb9837	informational	0.0	1.0	
89291	NaN	NaN	NaN	NaN	
113605	3f207df678b143eea3cee63160fa8bed	informational	0.0	1.0	
139992	3f207df678b143eea3cee63160fa8bed	informational	0.0	1.0	

	mobile	social	web	offer_id_redux_right	\
55972	1.0	1.0	0.0	NaN	
77705	1.0	1.0	0.0	5a8bc65990b245e5a138643cd4eb9837	
89291	NaN	NaN	NaN	5a8bc65990b245e5a138643cd4eb9837	
113605	1.0	0.0	1.0	NaN	
139992	1.0	0.0	1.0	3f207df678b143eea3cee63160fa8bed	

	offer_id
55972	5a8bc65990b245e5a138643cd4eb9837
77705	5a8bc65990b245e5a138643cd4eb9837
89291	5a8bc65990b245e5a138643cd4eb9837
113605	3f207df678b143eea3cee63160fa8bed
139992	3f207df678b143eea3cee63160fa8bed

```
[58]: with_full_offer_id = with_full_offer_id.drop(['offer_id_redux',
                                                    'offer_id_redux_right'],
                                                    axis = 1)
```

```
[59]: with_full_offer_id.columns
```

```
[59]: Index(['event', 'person', 'time', 'amount', 'reward_x', 'difficulty',
            'duration', 'id', 'offer_type', 'reward_y', 'email', 'mobile', 'social',
            'web', 'offer_id'],
            dtype='object')
```

Now we can merge it again to get the transaction events data. This means we'll duplicate many columns, but all we have to do is drop them.

```
[60]: with_transaction = with_full_offer_id.merge(processed_portfolio,
                                                    how = 'left',
                                                    left_on = 'offer_id',
                                                    right_on = 'id')
```

```
[61]: with_transaction.columns
```

```
[61]: Index(['event', 'person', 'time', 'amount', 'reward_x', 'difficulty_x',
            'duration_x', 'id_x', 'offer_type_x', 'reward_y', 'email_x', 'mobile_x',
            'social_x', 'web_x', 'offer_id', 'difficulty_y', 'duration_y', 'id_y',
```



```
'offer_type_y', 'reward', 'email_y', 'mobile_y', 'social_y', 'web_y'],
dtype='object')
```

```
[62]: with_transaction[['reward',
                        'reward_x',
                        'reward_y',
                        'difficulty_x',
                        'difficulty_y',
                        'duration_x',
                        'duration_y']]
```

```
[62]:
```

	reward	reward_x	reward_y	difficulty_x	difficulty_y	duration_x	\
0	0	NaN	0.0	0.0	0	3.0	
1	0	NaN	0.0	0.0	0	3.0	
2	0	NaN	NaN	NaN	0	NaN	
3	0	NaN	0.0	0.0	0	4.0	
4	0	NaN	0.0	0.0	0	4.0	
5	5	NaN	5.0	5.0	5	5.0	
6	0	NaN	NaN	NaN	0	NaN	
7	5	5.0	5.0	5.0	5	5.0	
8	5	NaN	5.0	5.0	5	5.0	
9	2	NaN	2.0	10.0	10	10.0	
10	5	NaN	NaN	NaN	5	NaN	
11	2	2.0	2.0	10.0	10	10.0	
12	2	NaN	2.0	10.0	10	10.0	
13	2	NaN	NaN	NaN	10	NaN	
14	2	NaN	2.0	10.0	10	7.0	
15	2	NaN	NaN	NaN	10	NaN	
16	2	2.0	2.0	10.0	10	7.0	
17	2	NaN	NaN	NaN	10	NaN	
18	2	NaN	NaN	NaN	10	NaN	
19	2	NaN	NaN	NaN	10	NaN	
20	5	NaN	5.0	5.0	5	5.0	
21	5	NaN	5.0	5.0	5	5.0	
22	5	NaN	NaN	NaN	5	NaN	
23	5	NaN	NaN	NaN	5	NaN	
24	5	NaN	NaN	NaN	5	NaN	
25	5	NaN	5.0	5.0	5	5.0	
26	5	NaN	5.0	5.0	5	5.0	
27	0	NaN	0.0	0.0	0	4.0	
28	0	NaN	0.0	0.0	0	4.0	
29	0	NaN	NaN	NaN	0	NaN	
...	
306504	2	2.0	2.0	10.0	10	10.0	
306505	2	NaN	NaN	NaN	10	NaN	
306506	2	NaN	NaN	NaN	10	NaN	
306507	5	NaN	5.0	20.0	20	10.0	

306508	5	NaN	5.0	20.0	20	10.0
306509	5	NaN	NaN	NaN	20	NaN
306510	5	5.0	5.0	20.0	20	10.0
306511	5	NaN	NaN	NaN	20	NaN
306512	5	NaN	NaN	NaN	20	NaN
306513	5	NaN	NaN	NaN	20	NaN
306514	5	NaN	NaN	NaN	20	NaN
306515	2	NaN	2.0	10.0	10	7.0
306516	2	NaN	2.0	10.0	10	7.0
306517	2	NaN	NaN	NaN	10	NaN
306518	2	2.0	2.0	10.0	10	7.0
306519	2	NaN	2.0	10.0	10	7.0
306520	2	NaN	2.0	10.0	10	7.0
306521	2	NaN	NaN	NaN	10	NaN
306522	2	2.0	2.0	10.0	10	7.0
306523	2	NaN	NaN	NaN	10	NaN
306524	5	NaN	5.0	5.0	5	7.0
306525	2	NaN	NaN	NaN	10	NaN
306526	5	5.0	5.0	5.0	5	7.0
306527	5	NaN	5.0	5.0	5	7.0
306528	2	NaN	2.0	10.0	10	7.0
306529	5	NaN	NaN	NaN	5	NaN
306530	2	2.0	2.0	10.0	10	7.0
306531	2	NaN	2.0	10.0	10	7.0
306532	2	NaN	NaN	NaN	10	NaN
306533	2	NaN	NaN	NaN	10	NaN

	duration_y
0	3
1	3
2	3
3	4
4	4
5	5
6	4
7	5
8	5
9	10
10	5
11	10
12	10
13	10
14	7
15	10
16	7
17	10
18	10

19	10
20	5
21	5
22	5
23	5
24	5
25	5
26	5
27	4
28	4
29	4
...	...
306504	10
306505	10
306506	10
306507	10
306508	10
306509	10
306510	10
306511	10
306512	10
306513	10
306514	10
306515	7
306516	7
306517	7
306518	7
306519	7
306520	7
306521	7
306522	7
306523	7
306524	7
306525	7
306526	7
306527	7
306528	7
306529	7
306530	7
306531	7
306532	7
306533	7

[306534 rows x 7 columns]

```
[63]: with_transaction(['reward',
                        'reward_x',
```

```

'reward_y',
'difficulty_x',
'difficulty_y',
'duration_x',
'duration_y']] .isnull().sum()

```

```

[63]: reward          0
      reward_x      272955
      reward_y      138953
      difficulty_x   138953
      difficulty_y      0
      duration_x     138953
      duration_y      0
      dtype: int64

```

We'll keep the `reward_y` column instead of the other ones because the lack of information there is related to the event type.

```

[64]: with_transaction.columns

```

```

[64]: Index(['event', 'person', 'time', 'amount', 'reward_x', 'difficulty_x',
           'duration_x', 'id_x', 'offer_type_x', 'reward_y', 'email_x', 'mobile_x',
           'social_x', 'web_x', 'offer_id', 'difficulty_y', 'duration_y', 'id_y',
           'offer_type_y', 'reward', 'email_y', 'mobile_y', 'social_y', 'web_y'],
          dtype='object')

```

```

[65]: with_transaction = with_transaction.drop(['reward',
                                              'reward_x',
                                              'difficulty_x',
                                              'duration_x',
                                              'id_x',
                                              'id_y',
                                              'offer_type_x',
                                              'email_x',
                                              'mobile_x',
                                              'social_x',
                                              'web_x'],
                                              axis = 1)

with_transaction.columns

```

```

[65]: Index(['event', 'person', 'time', 'amount', 'reward_y', 'offer_id',
           'difficulty_y', 'duration_y', 'offer_type_y', 'email_y', 'mobile_y',
           'social_y', 'web_y'],
          dtype='object')

```

```
[66]: with_transaction = with_transaction.rename(columns = {'reward_y': 'reward',
                                                         'difficulty_y': 'difficulty',
                                                         'duration_y': 'duration',
                                                         'offer_type_y': 'offer_type',
                                                         'email_y': 'email',
                                                         'mobile_y': 'mobile',
                                                         'social_y': 'social',
                                                         'web_y': 'web'})
```

```
[67]: with_transaction.head()
```

```
[67]:
```

	event	person	time	amount	reward	\
0	offer received	0009655768c64bdeb2e877511632db8f	168	NaN	0.0	
1	offer viewed	0009655768c64bdeb2e877511632db8f	192	NaN	0.0	
2	transaction	0009655768c64bdeb2e877511632db8f	228	22.16	NaN	
3	offer received	0009655768c64bdeb2e877511632db8f	336	NaN	0.0	
4	offer viewed	0009655768c64bdeb2e877511632db8f	372	NaN	0.0	

	offer_id	difficulty	duration	offer_type	\
0	5a8bc65990b245e5a138643cd4eb9837	0	3	informational	
1	5a8bc65990b245e5a138643cd4eb9837	0	3	informational	
2	5a8bc65990b245e5a138643cd4eb9837	0	3	informational	
3	3f207df678b143eea3cee63160fa8bed	0	4	informational	
4	3f207df678b143eea3cee63160fa8bed	0	4	informational	

	email	mobile	social	web
0	1	1	1	0
1	1	1	1	0
2	1	1	1	0
3	1	1	0	1
4	1	1	0	1

1.4 Finding successful offers

1.4.1 Defining success

To train the proposed models, we need to define success. For this problem, it makes sense that success is when a user has completed the funnels mentioned above successfully.

Let's start by finding which users got to which steps of the funnel. For that, we'll separate the events in its own tables and create auxiliary columns. Then we'll join those datasets to come up with a single dataset that has the offer id, the person id, and three booleans for received, viewed and completed. We'll also include the considerations for the order of the performed events to be considered valid in this step.

```
[68]: def creating_event_df(original_df, event):
    '''
    this function takes in a dataframe that contains the 'event',
    'person' and 'time' columns and creates a new dataframe for only
    one selected event sorted by people and time.

    INPUT
    - original_df (df): dataframe that contains the 'event',
      'person' and 'time' columns
    - event (str): the target event

    OUTPUT
    - target_df (df): sorted dataframe with the desired filters
    '''
    target_df = original_df[original_df['event'] == event].copy()
    target_df.sort_values(['person', 'time'])

    return target_df
```

```
[69]: offer_received = creating_event_df(with_transaction, 'offer received')

offer_received.rename(columns = {'time': 'time_received'},
                      inplace = True)

offer_received = offer_received.drop(['event',
                                     'amount',
                                     'difficulty',
                                     'offer_type',
                                     'email',
                                     'mobile',
                                     'social',
                                     'web'],
                                    axis = 1)

offer_received['received'] = np.ones(len(offer_received))
```

```
[70]: offer_viewed = creating_event_df(with_transaction, 'offer viewed')

offer_viewed.rename(columns = {'time': 'time_viewed'},
                    inplace = True)

offer_viewed = offer_viewed.drop(['event',
                                  'amount',
                                  'reward',
                                  'duration',
                                  'difficulty',
                                  'offer_type'],
```

```

        'email',
        'mobile',
        'social',
        'web'],
        axis = 1)

```

```
offer_viewed['viewed'] = np.ones(len(offer_viewed))
```

```
[71]: offer_transaction = creating_event_df(with_transaction, 'transaction')
```

```
offer_transaction.rename(columns = {'time': 'time_transaction'},
                        inplace = True)
```

```
offer_transaction = offer_transaction.drop(['event',
                                           'reward',
                                           'duration',
                                           'difficulty',
                                           'offer_type',
                                           'email',
                                           'mobile',
                                           'social',
                                           'web'],
                                           axis = 1)
```

```
offer_transaction['has_transaction'] = np.ones(len(offer_transaction))
```

```
[72]: offer_completed = creating_event_df(with_transaction, 'offer completed')
```

```
offer_completed.rename(columns = {'time': 'time_completed'},
                      inplace = True)
```

```
offer_completed = offer_completed.drop(['event',
                                       'amount',
                                       'reward',
                                       'duration',
                                       'difficulty',
                                       'offer_type',
                                       'email',
                                       'mobile',
                                       'social',
                                       'web'],
                                       axis = 1)
```

```
offer_completed['completed'] = np.ones(len(offer_completed))
```

```
[73]: # merge event dataframes into one
```

```
offers_received_viewed = offer_received.merge(offer_viewed,
```

```

                                how = 'left',
                                on = ['person', 'offer_id'])

offers_received_viewed_transaction = offers_received_viewed.
    ↳merge(offer_transaction,
                                how = 'left',
                                on =
    ↳['person', 'offer_id'])

offers = offers_received_viewed_transaction.merge(offer_completed,
                                how = 'left',
                                on = ['person', 'offer_id'])

# reorder columns
offers = offers[['person',
                  'offer_id',
                  'received',
                  'viewed',
                  'has_transaction',
                  'completed',
                  'duration',
                  'time_received',
                  'time_viewed',
                  'time_transaction',
                  'time_completed',
                  'amount',
                  'reward']]

# add expiration time for the offer
offers['time_expiry'] = offers['duration'] + offers['time_received']

```

```

[74]: # filter only pairs of events and person where the
# funnel order was followed correctly and the offer
# was successful
successful_offers = offers[(offers['has_transaction'] == 1) &
    ↳ # the offer must have had a transaction
                                (offers['time_transaction'] >=
    ↳offers['time_viewed']) & # the moment the transaction happened must be
    ↳after the offer was viewed
                                (offers['time_viewed'] >= offers['time_received']) &
    ↳ # the moment the offer was viewed must be after the offer was received
                                ((offers['time_expiry'] < offers['time_completed']))
    ↳| # the offers must have been completed - either by transaction in the
    ↳case

```



```

        (offers['time_expiry'] <
        →offers['time_transaction']))]    ## informational offers or by event -
        →before the expiration time

# add a column of ones signaling those are successful offers
successful_offers['is_successful'] = np.ones(len(successful_offers))

```

/home/julia.tessler/anaconda3/lib/python3.7/site-
packages/ipykernel_launcher.py:11: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>
This is added back by InteractiveShellApp.init_path()

```
[75]: successful_offers.shape
```

```
[75]: (217659, 15)
```

```

[76]: # filter only pairs of events and person where the
# funnel order was followed correctly but the offer
# was not successful
non_successful_offers = offers[offers['has_transaction'].isna()]    # if the
        →offer has no transaction, it wasn't successful

# add a column of ones signaling those are successful offers
non_successful_offers['is_successful'] = np.zeros(len(non_successful_offers))

```

/home/julia.tessler/anaconda3/lib/python3.7/site-
packages/ipykernel_launcher.py:7: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>
import sys

```
[77]: non_successful_offers.shape
```

```
[77]: (27573, 15)
```

Now we can merge it back into one dataset and then merge this data with the rest of the profile data.

```

[78]: offers = successful_offers.merge(non_successful_offers,
        how = 'outer')

```

```
[79]: offers.shape
```

```
[79]: (245232, 15)
```

We also need to refine user data. Let's start by joining the profile data with the transactions.

```
[80]: len(offers['person'].unique())
```

```
[80]: 16993
```

```
[81]: len(processed_profile['id'].unique())
```

```
[81]: 17000
```

```
[82]: len(with_transaction['person'].unique())
```

```
[82]: 17000
```

The amount of distinct users in our `offers` data isn't the same as it is in the `profile` data, which means there is a chance we have most of our offers covered by the `profile` data.

The `offers` dataset has 245,232 rows.

```
[83]: selected_users_profile = offers['person'].isin(processed_profile['id'])
      selected_users_offers = offers[selected_users_profile]
```

```
[84]: selected_users = with_transaction['person'].
      ↪isin(selected_users_offers['person'])
      selected_transactions = with_transaction[selected_users]
```

```
[85]: selected_transactions = selected_transactions.drop(['event',
                                                         'time'], axis = 1)
```

```
[86]: selected_transactions.shape
```

```
[86]: (306508, 11)
```

Now we have to filter this dataset only for the successful or non-successful offers we found above.

```
[87]: successful_transactions = offers.merge(selected_transactions,
                                             how = 'inner',
                                             right_on = ['person', 'offer_id'],
                                             left_on = ['person', 'offer_id'])
```

```
[88]: successful_transactions.columns
```

```
[88]: Index(['person', 'offer_id', 'received', 'viewed', 'has_transaction',
          'completed', 'duration_x', 'time_received', 'time_viewed',
```

```

'time_transaction', 'time_completed', 'amount_x', 'reward_x',
'time_expiry', 'is_successful', 'amount_y', 'reward_y', 'difficulty',
'duration_y', 'offer_type', 'email', 'mobile', 'social', 'web'],
dtype='object')

```

```

[89]: successful_transactions = successful_transactions.drop(['reward_x',
                                                            'received',
                                                            'viewed',
                                                            'has_transaction',
                                                            'completed',
                                                            'duration_x',
                                                            'time_received',
                                                            'time_viewed',
                                                            'time_transaction',
                                                            'time_completed',
                                                            'amount_x',
                                                            'time_expiry',
                                                            ],
                                                            axis = 1)

successful_transactions.rename(columns = {'reward_y': 'reward',
                                         'duration_y': 'duration',
                                         'amount_y': 'amount'},
                              inplace = True)

```

```

[90]: successful_transactions.shape

```

```

[90]: (2472021, 12)

```

Now we have multiple lines for the same successful transaction. Let's leave only one line of informations for each person and offer pair.

```

[91]: successful_transactions = successful_transactions.drop_duplicates()
      successful_transactions.head()

```

```

[91]:
      person                                     offer_id \
0  0009655768c64bdeb2e877511632db8f  5a8bc65990b245e5a138643cd4eb9837
2  0009655768c64bdeb2e877511632db8f  5a8bc65990b245e5a138643cd4eb9837
3  0009655768c64bdeb2e877511632db8f  3f207df678b143eea3cee63160fa8bed
5  0009655768c64bdeb2e877511632db8f  3f207df678b143eea3cee63160fa8bed
6  0009655768c64bdeb2e877511632db8f  f19421c1d4aa40978ebb69ca19b0e20d

      is_successful  amount  reward  difficulty  duration  offer_type  email \
0                1.0    NaN    0.0           0           3  informational    1
2                1.0   22.16    NaN           0           3  informational    1
3                1.0    NaN    0.0           0           4  informational    1
5                1.0    8.57    NaN           0           4  informational    1

```

6	1.0	NaN	5.0	5	5	bogo	1
	mobile	social	web				
0	1	1	0				
2	1	1	0				
3	1	0	1				
5	1	0	1				
6	1	1	1				

We still have some duplicates due to the multiple lines related to events... Since the **amount** column only happens when a user makes a transaction, this feature is highly important to predict the success because it's part of the definition of success. Therefore, it can't be used in the model and should be dropped.

```
[92]: transactions_logs = successful_transactions.drop(['amount'], axis = 1)

transactions_logs = transactions_logs.dropna().drop_duplicates()
```

```
[93]: transactions_logs.shape
```

```
[93]: (62398, 11)
```

Now we can join the user data with the transaction logs.

```
[94]: transactions_with_users = transactions_logs.merge(processed_profile,
                                                         how = 'left',
                                                         left_on = 'person',
                                                         right_on = 'id')
```

```
[95]: transactions_with_users.columns
```

```
[95]: Index(['person', 'offer_id', 'is_successful', 'reward', 'difficulty',
            'duration', 'offer_type', 'email', 'mobile', 'social', 'web', 'age',
            'gender', 'id', 'income', 'became_member_on_year'],
            dtype='object')
```

Cool! Now we're almost ready to start modelling. To model data, we need to remove the **person/id** and **offer_id** columns as well as encode the **offer_type** and **gender** (we'll use [Scikit-learn's LabelEncoder\(\)](#) for this). The amount coly

```
[96]: final_dataset = transactions_with_users.drop(['person',
                                                    'offer_id',
                                                    'id'],
                                                    axis = 1)
```

```
[97]: le = LabelEncoder()
```

```
[98]: offer_types_list = list(final_dataset['offer_type'].unique())
types_encoder = le.fit(offer_types_list)
offer_type_encoded = le.transform(final_dataset['offer_type'])

final_dataset['offer_type_encoded'] = offer_type_encoded
```

```
[99]: gender_list = list(final_dataset['gender'].unique())
gender_encoder = le.fit(gender_list)
gender_encoded = le.transform(final_dataset['gender'])

final_dataset['gender_encoded'] = gender_encoded
```

```
[100]: final_dataset = final_dataset.drop(['offer_type',
                                          'gender'], axis = 1)
```

```
[101]: final_dataset.columns
```

```
[101]: Index(['is_successful', 'reward', 'difficulty', 'duration', 'email', 'mobile',
            'social', 'web', 'age', 'income', 'became_member_on_year',
            'offer_type_encoded', 'gender_encoded'],
            dtype='object')
```

```
[102]: final_dataset.isna().sum()
```

```
[102]: is_successful      0
reward                  0
difficulty              0
duration               0
email                  0
mobile                 0
social                 0
web                   0
age                   0
income                0
became_member_on_year  0
offer_type_encoded     0
gender_encoded         0
dtype: int64
```

```
[103]: final_dataset.shape
```

```
[103]: (62398, 13)
```

```
[104]: final_dataset.head()
```

```
[104]:   is_successful  reward  difficulty  duration  email  mobile  social  web  \
0             1.0      0.0           0          3       1       1       1   0
```

1	1.0	0.0	0	4	1	1	0	1
2	1.0	5.0	5	5	1	1	1	1
3	1.0	2.0	10	10	1	1	1	1
4	1.0	5.0	5	5	1	1	1	1

	age	income	became_member_on_year	offer_type_encoded	\
0	33	72000.000000	2017	2	
1	33	72000.000000	2017	2	
2	33	72000.000000	2017	0	
3	33	72000.000000	2017	1	
4	118	65404.991568	2018	0	

	gender_encoded
0	1
1	1
2	1
3	1
4	3

1.5 Modeling

1.5.1 Separating data into train and test

We'll use [Scikit-learn's `train_test_split` function](#) to separate data, with 80% as train and 20% as test.

Our independent variables - or features - will be: * reward * difficulty * duration * email * mobile * social * web * amount * age * income * became_member_on_year * offer_type_encoded * gender_encoded

We'll train supervised models for this. Our dependent variable will be `is_successful`.

The supervised models that will be trained are: * Model 1: Logistic Regression * Model 2: Naïve Bayes * Model 3: Support Vector Machines (SVM) * Model 4: Decision Tree

We'll use Scikit-learn's framework for all models.

```
[105]: features = final_dataset[['reward',
                                'difficulty',
                                'duration',
                                'email',
                                'mobile',
                                'social',
                                'web',
                                'age',
                                'income',
                                'became_member_on_year',
                                'offer_type_encoded',
                                'gender_encoded']]
```

```
y = final_dataset['is_successful']
```

```
[106]: X_train, X_test, y_train, y_test = train_test_split(features,
                                                         y,
                                                         test_size = 0.8,
                                                         random_state = 42)
```

To train all models, we'll use [Scikit-learn's cross_val_score](#) with 5 folds.

1.5.2 Baseline model

As baseline model we'll use a naive one, given by [Scikit-learn's DummyClassifier](#). The model shall predict by most frequent class. This is a great baseline model because it's no better than just saying that an offer will be a success just because that's the majority's class (and, since this is an unbalanced problem, this means that predictions by the majority will be better than flipping a fair coin to predict the class).

```
[107]: dummy_clf = DummyClassifier(strategy = "most_frequent")

cross_val_baseline = cross_val_score(dummy_clf,
                                     X_train,
                                     y_train,
                                     cv = 5,
                                     scoring = 'f1_micro')
```

```
[108]: cross_val_baseline.mean()
```

```
[108]: 0.6441221316975467
```

The baseline model has 64.41% of F1-Score. We're looking for models that can perform better than this one.

1.5.3 Model 1: Logistic Regression

The [logistic regression](#) is a fairly simple model that presents good results in robust datasets, such as this one. We expect it to be better than the baseline and, since this is one of the most explainable models there is, we might favor it depending on the results. The logistic regression (along with the Decision Tree) is one of the models with the easiest explanation we chose to train, which makes it a favourite to be the best model (but first let's see how the models perform).

Even though the logistic regression has analytical solution, Scikit-learn's implementation uses computational methods to solve it. Since this is a small dataset, we'll use the `liblinear` solver as suggested by [the documentation](#).

```
[109]: lr = LogisticRegression(solver = 'liblinear', random_state = 42)

cross_val_lr = cross_val_score(lr,
                               X_train,
                               y_train,
```

```
cv = 5,  
scoring = 'f1_micro')
```

```
[110]: cross_val_lr.mean()
```

```
[110]: 0.6662358493158751
```

We can see some marginal gain in the F1-Score: this model reached 66.62%.

1.5.4 Model 2: Naïve Bayes

We'll train a Gaussian Naïve Bayes model. This model usually performs well in well defined classification models. The chosen classifier, Gaussian, assumes that the likelihood of the features comes from a Gaussian distribution (see [documentation](#) for more details).

Naïve Bayes classifiers use [Bayesian statistics](#) to find the parameter values. An advantage of Naïve Bayes is that it only requires a small number of training data to estimate the parameters necessary for classification, even though this is a naïve model, which makes it great for this problem. Also, it can perform very well for this type of problem, even when the data breaks the Gaussian assumption.

```
[111]: nb = GaussianNB()  
  
cross_val_nb = cross_val_score(nb,  
                                X_train,  
                                y_train,  
                                cv = 5,  
                                scoring = 'f1_micro')
```

```
[112]: cross_val_nb.mean()
```

```
[112]: 0.702621055986057
```

This model performed a bit better than the Logistic Regression and baseline models when it comes to F1-Score: 70.26%.

1.5.5 Model 3: Support Vector Machines (SVM)

We'll train support vector machines. SVMs are considered more complex than the models above. We expect it to perform better than the baseline, but it might start to overfit the model.

An SVM model builds a representation of the categories as points in space, mapped in a way that the categories are each divided by a clear gap that is as wide as possible (see more [here](#)). Also, SVMs can efficiently perform a non-linear classification using what is called the kernel trick, implicitly mapping their inputs into high-dimensional feature spaces.

For this problem, we'll use an [SVC kernel](#), which is a linear one.

```
[113]: sv = svm.SVC()  
  
cross_val_sv = cross_val_score(sv,
```



```
X_train,  
y_train,  
cv = 5,  
scoring = 'f1_micro')
```

```
/home/julia.tessler/anaconda3/lib/python3.7/site-  
packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will  
change from 'auto' to 'scale' in version 0.22 to account better for unscaled  
features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.
```

```
"avoid this warning.", FutureWarning)
```

```
/home/julia.tessler/anaconda3/lib/python3.7/site-  
packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will  
change from 'auto' to 'scale' in version 0.22 to account better for unscaled  
features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.
```

```
"avoid this warning.", FutureWarning)
```

```
/home/julia.tessler/anaconda3/lib/python3.7/site-  
packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will  
change from 'auto' to 'scale' in version 0.22 to account better for unscaled  
features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.
```

```
"avoid this warning.", FutureWarning)
```

```
/home/julia.tessler/anaconda3/lib/python3.7/site-  
packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will  
change from 'auto' to 'scale' in version 0.22 to account better for unscaled  
features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.
```

```
"avoid this warning.", FutureWarning)
```

```
/home/julia.tessler/anaconda3/lib/python3.7/site-  
packages/sklearn/svm/base.py:193: FutureWarning: The default value of gamma will  
change from 'auto' to 'scale' in version 0.22 to account better for unscaled  
features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.
```

```
"avoid this warning.", FutureWarning)
```

```
[114]: cross_val_sv.mean()
```

```
[114]: 0.6510127408915569
```

This model has 65.10% of F1-Score. It is better than the baseline model, but not better than the Naïve Bayes model.

1.5.6 Model 4: Decision Tree

We'll train a decision tree classifier. This is the most powerful of all the chosen models and is very likely to overfit.

Decision trees are highly explainable because they produce a [flow-chart like structure](#) as outcome, which means that we have a linear flow of decision making in order to classify the offers. It's a favourite to best model due to it's explainability (as mentioned above).

```
[115]: dt = DecisionTreeClassifier(random_state = 42)

cross_val_dt = cross_val_score(dt,
                                X_train,
                                y_train,
                                cv = 5,
                                scoring = 'f1_micro')
```

```
[116]: cross_val_dt.mean()
```

```
[116]: 0.6508533831240313
```

This better than the baseline model, since the F1-Score is 65.08%. But it's not the best. The Linear Regression model has better scores and it's explainable as well.

1.6 Choosing the best model

Based on the results we found above, the best model is the Naïve Bayes. Let's retrain it so we can keep it's coefficients and apply it to the test features.

```
[117]: model = nb.fit(X_train, y_train)
```

```
[118]: predictions = model.predict(X_test)
```

```
[119]: # get the confusion matrix
pd.crosstab(y_test, predictions)
```

```
[119]: col_0      0.0    1.0
is_successful
0.0          6095  12012
1.0          3097  28715
```

```
[120]: metrics.f1_score(y_test, predictions, average = 'micro')
```

```
[120]: 0.6973296740719966
```

```
[121]: print(metrics.classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0.0	0.66	0.34	0.45	18107
1.0	0.71	0.90	0.79	31812
accuracy			0.70	49919
macro avg	0.68	0.62	0.62	49919
weighted avg	0.69	0.70	0.67	49919

The model performed not very well in the test dataset, reaching a 69.73% F1-Score. The decrease in the score was already expected, since this data is new to the model. But since this score is close to what we found during the train stage, this means that this model can be applied to other datapoints and still be trustworthy, as it'll produce similar results.

The other usual metrics performed fairly well, with exception of the recall for non-successful offers, which was 34% - meaning that it's prediction is worse than flipping a coin to predict the success of the offer.

Unfortunately, Naïve Bayes model do not have easy explanations for the coefficients, which means that we can't easily explain what variables are related to the success of the offer. What we can do is look into each variable and the associated probability to look into the relations.

```
[122]: result_df = pd.DataFrame(X_train.columns)
result_df[['non successful', 'successful']] = pd.DataFrame(model.
    ↳predict_proba(X_train).tolist())

result_df
```

```
[122]:
```

	0	non successful	successful
0	reward	0.204731	0.795269
1	difficulty	0.925091	0.074909
2	duration	0.346706	0.653294
3	email	0.294267	0.705733
4	mobile	0.408526	0.591474
5	social	0.093768	0.906232
6	web	0.924745	0.075255
7	age	0.337013	0.662987
8	income	0.398634	0.601366
9	became_member_on_year	0.288445	0.711555
10	offer_type_encoded	0.118418	0.881582
11	gender_encoded	0.086060	0.913940

Also we can get the features that had the most predictive contribution to each class, but still no explainability.

```
[123]: # source: https://stackoverflow.com/questions/50526898/
    ↳how-to-get-feature-importance-in-naive-bayes
# prints the top 10 most predictive features for the non-successful class
neg = model.theta_[0].argsort()
print(np.take(X_train.columns, neg[:10]))

print('')

# prints the top 10 most predictive features for the successful class
neg = model.sigma_[0].argsort()
print(np.take(X_train.columns, neg[:10]))
```

```
Index(['social', 'mobile', 'web', 'gender_encoded', 'offer_type_encoded',
```

```
    'email', 'reward', 'duration', 'difficulty', 'age'],
    dtype='object')
```

```
Index(['email', 'web', 'mobile', 'social', 'offer_type_encoded',
      'gender_encoded', 'became_member_on_year', 'duration', 'reward',
      'difficulty'],
      dtype='object')
```

If we consider the Logist Regression as the best explainable model, we can look into the coefficients to get the most important features.

```
[124]: model_lr = lr.fit(X_train, y_train)
```

```
[125]: predictions_lr = model_lr.predict(X_test)
```

```
[126]: metrics.f1_score(y_test, predictions_lr, average = 'micro')
```

```
[126]: 0.6438029607964904
```

```
[127]: # get the confusion matrix
pd.crosstab(y_test, predictions_lr)
```

```
[127]: col_0      0.0    1.0
is_successful
0.0          523   17584
1.0          197   31615
```

```
[128]: # Logistic Regression needs a small transformation to get the right coefficients
transformed_coefficients = list(np.exp(model_lr.coef_))
cdf = pd.DataFrame(list(X_train.columns), transformed_coefficients).
    ↪reset_index()

cdf.columns = ['coefficient', 'feature']
print(cdf.sort_values(by = 'coefficient', ascending = False))
```

	coefficient	feature
0	1.051222	reward
5	1.023997	social
4	1.008203	mobile
9	1.000550	became_member_on_year
3	1.000011	email
8	0.999996	income
11	0.999768	gender_encoded
7	0.998785	age
6	0.995864	web
2	0.992067	duration
10	0.987936	offer_type_encoded
1	0.960984	difficulty

The supervised models trained, with respective F1-Scores, were: * Baseline Model: Dummy Classifier for most frequent class * F1-Score: 64.41% * Model 1: Logistic Regression * F1-Score: 66.62% * Model 2: Naïve Bayes * F1-Score: 70.26% * Model 3: Support Vector Machines (SVM) * 65.10% * Model 4: Decision Tree * F1-Score: 65.08%

The best model was Naïve Bayes, but it has not explainability. The second best and explainable model was the Logistic Regression.

In order to have explainability, we loose predictive power: the F1-Score of the Logistic Regression model is 64.38%. But the value of the reward, followed by the offer sent by social networks. For the increase of one unit in each feature, we expect a increase of the respective coefficient in the success of the offer.

1.7 Conclusion and refinement

When we started, we wanted to make better purchasing offers to Starbucks' customers. For this, we used customer's past behaviour to find patterns and try to be more assertive. As given by the Udacity's Starbucks Project Overview, the basic task was to use the data to identify which groups of people are most responsive to each type of offer, and how best to present each type of offer. In other words, this is a classification problem where the model takes user behaviour data as input and produces a group as output (either previously defined or not).

For this project, we spent quite some time dealing with the features, manipulating those to fit into the models. For that to happen, we found a way to define an offer success based on the user funnel performed from the transcript dataset.

Once we had the dataset, we trained 4 supervised learning models and a baseline one. The baseline was an incredibly naïve model that classified the items based on the most frequent class. The model with best performance was a Naïve Bayes. This model doesn't have easy explainability, which means we fail to find understandable patterns to provide offers. But this model achieved reasonable results when applied to the test dataset, meaning it can be applied to other datapoints and still produce the same level of results.

In order to get explainability, we chose the Logistic Regression model, that has lower predictive power. But with this model, we identified the top three most important features: reward, social and mobile.

Next steps would include better feature engineering and selection (we just used all features we could) and other classification models. The model selection should probably account for model explainability, which failed in this case. To make better decisions here, we could dive deeper into the post-mortem analysis of those models.

[]: